

WG01: Kernel C++

Amlal El Mahrouss
NeKernel.org

December 6, 2025

1 Abstract

Many and most kernels have been shipped using the C programming language.

And some of them like EKA2 uses the C++ programming language. Although notoriously difficult, one may still adapt to those constraints in order to deliver one such operating system kernel.

That is the reason that most production-grade kernels (Linux, XNU, and NT) are mostly written in C. With a higher-level subset in C++.

However, when correctly applying C++ principles to kernel development, one makes the development much more easier to pull off.

2 The Three Principles of Kernel C++

2.1 Part One: The C++ Runtime

The problem mostly lies in the C++ runtime. Which assumes an existing host. A host is the contrary of a freestanding target, that is a program which expects a runtime to be present and linked to the program.

A C++ Kernel may instead make use of compile-time features of C++ alongside a tiny C++ runtime to make sure that no issues arise because of this host/freestanding difference.

2.2 Part Two: Constexpr and Friends

One may avoid V-tables or runtime dependent features when possible. While focusing instead on meta-programming and compile-time features offered by C++. For example one may use templates to implement a scheduling policy algorithm.

One example of such implementation may be:

(Keep in mind that the code has not been tested on production systems, use it at your own risk!)

```

1  /// Reference Implementation: https://github.com
2  /// /nekernel-org/nekernel/blob/stable/src/kernel/
3  /// KernelKit/CoreProcessScheduler.h#L51-L76
4  /// AND: https://github.com/nekernel-org/nekernel/blob/stable/
5  /// src/kernel
6  /// /KernelKit/CoreProcessScheduler.h#L78-L105
7
8  struct FileTree final {
9      static constexpr bool is_virtual_memory = false;
10     static constexpr bool is_memory = false;
11     static constexpr bool is_file = true;
12     /// ...
13 };
14
15 struct MemoryTree final {
16     static constexpr bool is_virtual_memory = false;
17     static constexpr bool is_memory = true;
18     static constexpr bool is_file = false;
19     /// ...
20 };

```

As you see, two structures leverages the ‘constexpr’ keyword to make sure no bugs or panic occur at runtime because of a misuse of a system resource.

Which is why the constexpr keyword is very powerful here, we avoid the many pitfalls of writing (and hoping) that the C version will be well-thought enough so that we can catch such bugs later.

2.3 Bonus: Using Concepts and the DDK.

This bonus subsection intends to show how properly applied C++ Can solve issues related to driver validation. The following example shows how powerful C++ can be when combining it with Device Driver development as well.

```

1  /// Reference implementation:
2  /// https://github.com/nekernel-org
3  /// /nekernel/blob/stable/src/libDDK/DriverKit
4  /// /c%2B%2B/driver_base.h
5  /// @brief This concept requires the driver
6  /// to be IDriverBase compliant.
7
8  constexpr auto kInvalidType = 0;
9
10 template <typename T>
11 concept IsValidDriver = requires(T a) {
12     { a.IsActive() && a.Type() > kInvalidType };
13 };
14
15 /// @brief Consteval helper to detect

```

```

16 /// whether a template is truly based on IDriverBase.
17 /// @note This helper is consteval only.
18 template<IsValidDriver T>
19 consteval void ce_ddk_is_valid(T) {}

```

2.4 Part Three: Memory and C++ Classes

This last part treats about the final and most important part of this paper so far. Memory. As you may already have known, the C++ language uses a class lookup system (also called v-table) in order to refer to a base method in case if the instance has it missing.

```

1 /// Link: https://godbolt.org/z/q3Wceannr
2 /// /std:c++20 /Wall
3
4 #include <iostream>
5
6 class A
7 {
8 public:
9     explicit A() = default;
10    virtual ~A() = default;
11
12    virtual void doImpl()
13    {
14        std::cout << "doImpl\n";
15    }
16 };
17
18 class B : public A
19 {
20 public:
21     explicit B() = default;
22     ~B() override = default;
23 };
24
25 int main() {
26     B* callImpl = new B();
27     callImpl->doImpl();
28 }

```

3 Addressing V-Tables in a C++ Kernel.

Now the problem with kernel development is that we want to avoid such feature as much as possible, and we'd do that by following the Prong on Inheritance:

3.1 The Three Prongs on Inheritance:

Consider the following questions:

A: Is this a protocol/concept that can be extended to other similar protocols/concepts?

B: Can I do this without too much trade-off costs?

C: Can I do this without using V-Tables?

When 2/3 of those questions fail, you should consider finding another solution to your problem, as it surely has an equivalent without V-Tables.

4 Conclusion

We can now conclude that C++ in Kernel Development is indeed possible under strict conditions of C++ Development.

Breaking those conditions would lead to system quirks and instability. A reference implementation of this paper exists, It's called NeKernel.org, available under the same internet address.

I am looking forward to any questions or inquiries at: amlal@nekernel.org.

5 References

EKA2: [Link](#)

NeKernel.org: [Link](#)

Virtual Method Tables: [Link](#)